

COE347 HW4

Author: Alex Choi

Date: 2025-02-17

Question 1

The Forward Difference Approximation is:

$$u(x+h) = u(x) + hu'(x) + \frac{h^2}{2!}u''(x) + \frac{h^3}{3!}u'''(x) + \frac{h^4}{4!}u^{(4)}(x) + \dots$$

The Backward Differene Approximation is:

$$u(x-h) = u(x) - hu'(x) + \frac{h^2}{2!}u''(x) - \frac{h^3}{3!}u'''(x) + \frac{h^4}{4!}u^{(4)}(x) + \dots$$

Switch to index notation and add $u(x+h)$ and $u(x-h)$

$$u_{i+1} + u_{i-1} = 2u_i + h^2u''_i + 2\frac{h^4}{4!}u^{(4)}(x) + h.o.t.$$

Now, rearrange:

$$h^2u''_i = u_{i+1} - 2u_i + u_{i-1} - 2\frac{h^4}{4!}u^{(4)}(x_i) + h.o.t.$$

Here, we can clearly see that:

$$C = -2\frac{1}{4!} = -\frac{1}{12}$$

In [397...

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd # for tabular displays
from scipy.linalg import solve, solve_banded # for efficeint solver of tridiagonal system.
from scipy.interpolate import PchipInterpolator # need for 3rd order interpolation of exact solution
from scipy.stats import linregress
from pprint import pprint
from IPython.display import display, Markdown, HTML

# we will go ahead and define a helper function to read in exact solution file for latr use
def read_dat_file(filename = str) -> tuple:
    # Load the data from the file
    data = np.loadtxt(filename)
```

```
# Extract x and y values (assuming two columns)
return np.array(data[:, 0]), np.array(data[:, 1])

# let's first define our ODE, or a given function
def ode_eq1(x):
    """
    Returns f(x) = (4π)^2 cos(4πx)
    """
    return 16 * (np.pi ** 2) * np.cos(4 * np.pi * x)
```

Question 2

Here, we implement a numerical solution to the given boundary value problem using the second-order finite difference formula (FDF) for approximating $u''(x)$. We discretize the domain, construct a system of linear equations, and solve for u_i at interior points. The numerical solution is obtained for $N = 10$ and can be compared against the exact solution provided in the dataset.

We first have to attain a matrix-vector form of our linear systems of equations for all values of u_i

Given our ODE:

$$u''(x) = f(x) = (4\pi)^2 \cos(4\pi x) = 16\pi^2 \cos(4\pi x) \quad (1)$$

And Given our Second-Order Finite Difference Approximation:

$$u''(x_i)h^2 = u_{i+1} - 2u_i + u_{i-1} + \mathcal{O}(h^2) \quad (2)$$

Equating (1) and (2):

$$h^2 16\pi^2 \cos(4\pi x_i) = u_{i+1} - 2u_i + u_{i-1} \quad (3)$$

Converting to a System of Equations:

$$h^2 \begin{bmatrix} f_1 - u(0) \\ f_2 \\ \vdots \\ f_N - u(1) \end{bmatrix} = \begin{bmatrix} -2 & 1 & 0 & \dots & 0 \\ 1 & -2 & 1 & \dots & 0 \\ 0 & 1 & -2 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & 1 & -2 \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \\ \vdots \\ u_N \end{bmatrix}$$

where:

- The **main diagonal** is filled with **-2**.
- The **superdiagonal** and **subdiagonal** are filled with **1**.
- The right-hand side includes **boundary conditions** at $x = 0$ and $x = 1$.

This forms a **tridiagonal system** that we can solve numerically.

Construct the BVP solver

In [398...

```
# now formulate the BVP solver for our ODE:
def solve_bvp(ode, x_range: list, u_i: float, u_f: float, N: int):
    """
    Solves a boundary value problem (BVP) using the second-order finite difference method (FDF).

    Parameters:
        ode (function): The function representing the right-hand side of the ODE,  $f(x)$  in  $u''(x) = f(x)$ .
        u0 (float or int): The boundary condition at  $x = 0$  (i.e.,  $u(0) = u_0$ ).
        uf (float or int): The boundary condition at  $x = 1$  (i.e.,  $u(1) = u_f$ ).
        N (int): The number of interior grid points.

    Returns:
        numpy.ndarray: The numerical solution  $u$  at the interior grid points plus boundary conditions.
        numpy.ndarray: The corresponding  $x$  values at the interior points.

    Notes:
        - The step size  $h$  is computed as  $h = 1 / (N + 1)$ .
        - The finite difference approximation is used to construct a system of linear equations.
        - The system is solved to obtain the numerical solution for  $u(x)$  at interior points.
    """
    # Extract final and initial x values
    x_i, x_f = x_range[0], x_range[1]
    # define step-size h:
    h = (x_f - x_i) / (N + 1)

    # construct our tri-diagonal matrix system
    main_diag = -2 * np.ones(N, dtype = float) # main diagonal is -2s.
    other_diag = np.ones(N-1, dtype = float) # super and sub diagonal are all 1s. Note: subtract one row.
    # "stack" the diagonals in the format expected by `solve_banded`. ab has to be banded, not normal numpy matrix
    tri_diagonal = np.zeros((3, N))
    tri_diagonal[0, 1:] = other_diag # Superdiagonal (shifted right)
    tri_diagonal[1, :] = main_diag # Main diagonal
    tri_diagonal[2, :-1] = other_diag # Subdiagonal (shifted left)

    # define N number of interior points, x_i for i = 1,...,N
    x_interior = np.linspace(start = x_i + h, stop = x_f - h, num = N, dtype = float)

    # evaluate values of ODE, f(x_i), at different interior points
    f = ode(x_interior) * (h ** 2) # multiply by h^2
    # Apply boundary conditions to the right-hand side vector
    f[0] -= u_i # apply boundary condition u(0) = u_0
    f[-1] -= u_f # apply boundary condition u(f) = u_{N+1}

    # solve the system of equations
    u_interior = solve_banded(l_and_u = (1, 1), ab = tri_diagonal, b = f)

    # finally, form full x and u by including boundary conditions and interior values
    # create empty array to store interior plus two end points
    x = np.empty(N + 2)
    u = np.empty(N + 2)
    # assign boundary conditions
    x[0], x[-1] = x_i, x_f # starting and ending x values
    u[0], u[-1] = u_i, u_f # starting and ending u values
    # Assign interior values
    x[1:-1] = x_interior
    u[1:-1] = u_interior
```

```
return x, u
```

Plot of Numerical against Exact Solution for $N = 10^4$.

```
In [399... sns.set_style("darkgrid")
sns.set_palette("flare")

# problem conditions
N = 10
x_span = [0,1]
u0 = 0
u1 = 2


# load in the exact values
exact_x, exact_u = read_dat_file("solutionA_N10000.dat")
# comput data points
numerical_x_Ne4, numerical_u_Ne4 = solve_bvp(ode_eq1, x_span, u0, u1, N)

# Plot the data
plt.figure(figsize=(16, 6))
# Exact solution
sns.lineplot(x=exact_x, y=exact_u, marker='o', markersize=6, linestyle='-', linewidth=2.5, alpha=0.9, markevery=100, label="Exact")
# Numerical solution
sns.lineplot(x=numerical_x_Ne4, y=numerical_u_Ne4, marker='x', markersize=6, linestyle='--', linewidth=2.5, alpha=0.9, markevery=100, label="Numerical")

# xlabel, ylabel, title, etc.
plt.xlabel("x", fontsize=14), plt.ylabel("u(x)", fontsize=14), plt.title("Comparison of Exact and Numerical Solution for ODE, eq. (1)", fontsize=16)
plt.legend(fontsize=12), plt.grid(True, linestyle="--", alpha=1)

# Show plot
plt.show()
```

Comparison of Exact and Numerical Solution for ODE, eq. (1)



Solution for $N = 10$ and Tabular Format

We will compute our numerical solution for the case $N = 10$ and print the results in a tabular format.

```
In [400... N = 10
# obtain numerical solution for N = 10
numeric_x_N10, numeric_u_N10 = solve_bvp(ode_eq1, x_span, u_i=u0, u_f=u1, N=N)
# make a pandas DataFrame for tabular format
df_N10 = pd.DataFrame(
    data={
        '$x_i$': numeric_x_N10, # LaTeX notation for x_i
        '$u_i$': numeric_u_N10  # LaTeX notation for u_i
    },
    index=pd.Index(range(N+2), name='$N$') # Ensure index is also formatted in LaTeX
)

# Display DataFrame with LaTeX formatting in HTML
display(HTML(df_N10.to_html(escape=False)))
```

N		x_i	u_i
	0	0.000000	0.000000
	1	0.090909	0.834354
	2	0.181818	2.210854
	3	0.272727	2.732715
	4	0.363636	2.002368
	5	0.454545	1.086290
	6	0.545455	1.268108
	7	0.636364	2.547822
	8	0.727273	3.641805
	9	0.818182	3.483582
	10	0.909091	2.470718
	11	1.000000	2.000000

Question 3

Error Analysis for various values of N

We will compute the numerical solution for different values of N and compare it against the exact solution ($N = 10^4$). The errors are defined as:

$$E = \left[\sum_{i=1}^N (u_i - \tilde{u}_i)^2 \right]^{\frac{1}{2}}$$

$$e = \frac{1}{N} \left[\sum_{i=1}^N (u_i - \tilde{u}_i)^2 \right]^{\frac{1}{2}}$$

where:

- u_i is the numerical solution
- \tilde{u}_i is the interpolated exact solution at x_i
- $h = (N + 1)^{-1}$ is the step size.

We will compute the errors for:

$$N = \{5, 10, 20, 40, 80, 160, 320, 640, 1280\}$$

and fit a function of the form:

$$E = Ch^\alpha$$

Interpolation of the Exact Solution

Since the exact solution is only available at a fine grid ($N = 10^4$), we must interpolate it to match the grid points of the numerical solutions ($N = 5, 10, 20, \dots, 1280$).

To ensure accurate error calculations:

- The exact solution \tilde{u}_i should be evaluated **at the same x_i points** as the numerical solution.
- We use **shape-preserving interpolation** (`PchipInterpolator`) to minimize oscillations and avoid large interpolation errors.
- The interpolation error **must be smaller** than the numerical error we are analyzing.

```
In [401... # store different number of steps
steps = np.array([5, 10, 20, 40, 80, 160, 320, 640, 1280], dtype = int)
# we will first compute different step-sizes h and store it
step_sizes = 1 / (steps + 1)

# Solve BVP for multiple step values and store results in lists
numeric_solutions = [solve_bvp(ode_eq1, x_span, u0, u1, N) for N in steps]
# Extract x and u values for each N
numeric_Xs = [np.array(sol[0], dtype=float) for sol in numeric_solutions] # Extract all x values
numeric_Us = [np.array(sol[1], dtype=float) for sol in numeric_solutions] # Extract all u values
# Create an interpolator object for multiple uses
interpolator = PchipInterpolator(exact_x, exact_u)
# Interpolate "exact" solutions at different N values, INCLUDING boundary conditions
interpolated_exact_Us = [interpolator(Xs) for Xs in numeric_Xs]
```

Log-log plot that shows E and e versus h^{-1}

```
In [402... def compute_errors(N, exact: np.array, numerical: np.array) -> float:
    """
    Computes global and local errors.

    Parameters:
        N (int): Number of interior points
        exact (np.array): Interpolated exact solution at grid points
        numerical (np.array): Numerical solution

    Returns:
        float, float: Global error (E) and Local error (e)
    """
    # compute global error, this is essentially L2 norm, which we can use numpy's vectorized function.
    global_error = np.linalg.norm(exact - numerical)
    # compute the local error, scaling by N
    local_error = global_error / float(N)

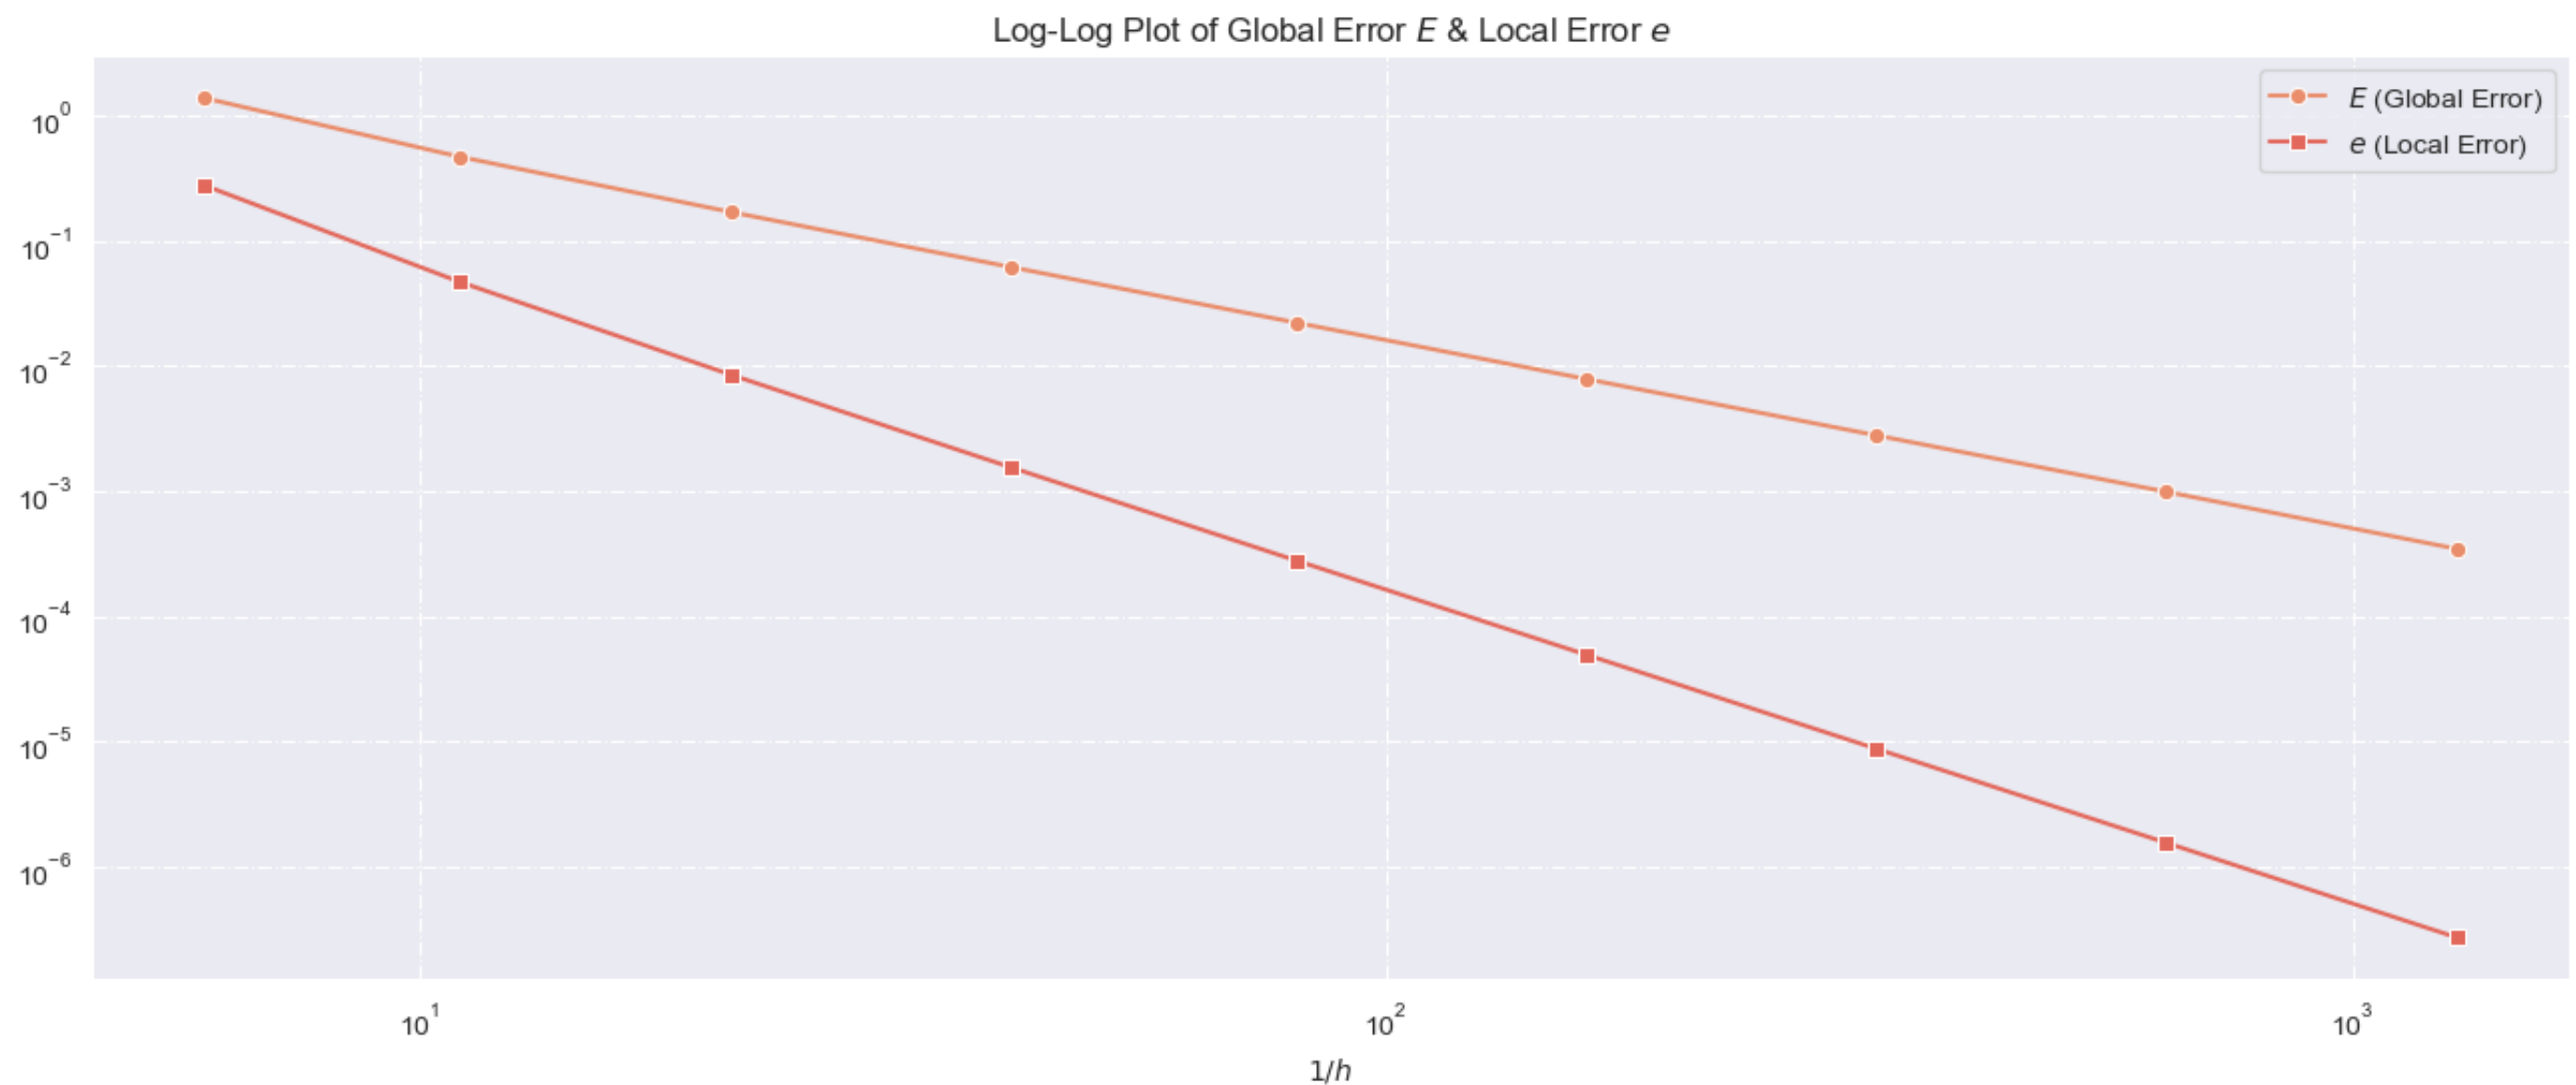
    return global_error, local_error

# compute errors for varying steps
global_errors = [compute_errors(steps[i], interpolated_exact_Us[i], numeric_Us[i])[0] for i in range(len(numeric_Us))] # global errors
local_errors = [compute_errors(steps[i], interpolated_exact_Us[i], numeric_Us[i])[1] for i in range(len(numeric_Us))] # local errors

# create a plot for E and e
plt.figure(figsize=(16, 6)) # create a subplot
sns.lineplot(x = 1 / step_sizes, y=global_errors, label=r"$E$ (Global Error)",marker='o')
```

```
# plot local error
sns.lineplot(x = 1 / step_sizes, y=local_errors, label=r"$e$ (Local Error)",marker='s')

# configurations
y_labels = ["Global Error $E$", "Local Error $e$"]
titles = ["Log-Log Plot of Global Error $E$", "Log-Log Plot of Local Error $e$"]
# apply all the configs
plt.xlabel(r"$1/h$"), plt.xscale('log'), plt.yscale('log')
plt.title("Log-Log Plot of Global Error $E$ & Local Error $e$") # set the tile
plt.legend() # show the legend
plt.grid(True, linestyle="dashdot", alpha=1.0)
plt.show()
```



Fitting Ch^α to Global and Local Errors

Here, we fit a function of the form Ch^α and report the value of α for both $\log(E)$ and $\log(e)$. Then, we make a conclusion about the order of the errors E and e with respect to $h = (N + 1)^{-1}$. Since we take the log of the function Ch^α , our function becomes:

$$\log(E) = \log(C) + \alpha \log(h)$$

$$\log(e) = \log(C) + \alpha \log(h)$$

This transformation allows us to express the relationship in a linear form:

$$y = mx + b$$

where $y = \log(E)$ (or $\log(e)$), $x = \log(h)$, slope $m = \alpha$, and intercept $b = \log(C)$.

```
In [403]: # For each global and local error, fit a linear function to get alpha
global_alpha = linregress(np.log(step_sizes), np.log(global_errors))[0] # global error alpha fitting
local_alpha = linregress(np.log(step_sizes), np.log(local_errors))[0] # local error alpha fitting

# Create a DataFrame with correct column structure
df_Q3_alphas = pd.DataFrame({
    'Type': ['$\alpha_E$', '$\alpha_e$'],
    'Value': [global_alpha, local_alpha]
})

# Display the table with formatted LaTeX notation for HTML rendering
display(HTML(df_Q3_alphas.round(3).to_html(index=False, escape=False)))
```

Type	Value
α_E	1.531
α_e	2.558

Interpretation of Results

The expected theoretical behavior for a **second-order** finite difference method is:

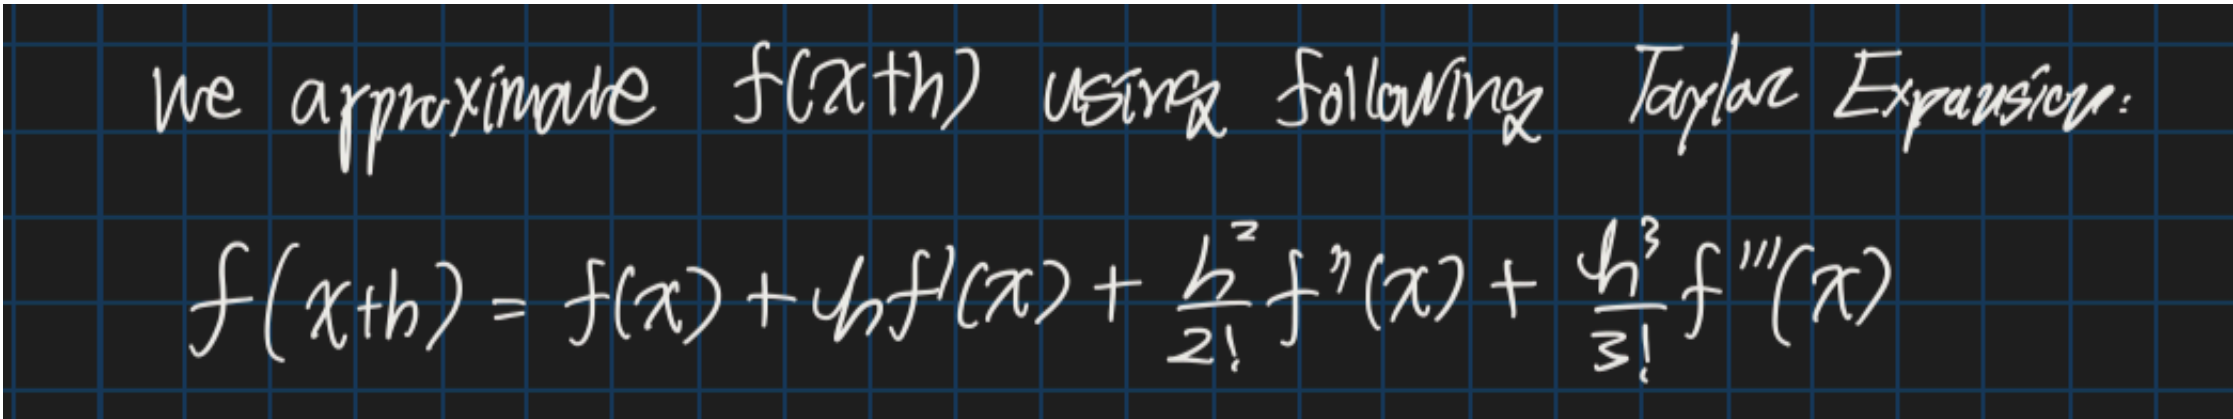
- $E = \mathcal{O}(h^2)$ (global error)
- $e = \mathcal{O}(h^3)$ (local error)

Our computed values of $\alpha_{Global} \approx 1.53$ and $\alpha_{Local} \approx 2.56$ show that:

- **The method is converging**, but the **global order of accuracy is slightly lower than expected (same for local)**.
- Possible reasons for this include **discretization effects at boundaries**, **floating-point precision errors**, or **not using a fine enough grid to see asymptotic behavior**.

Conclusion: The numerical method exhibits approximately second-order global convergence and higher-order local convergence. However, to verify full $\mathcal{O}(h^2)$ accuracy, additional refinements such as increasing N should be considered. Since the smallest h is only approximately 0.000781, floating-point precision errors should not be a dominant factor at this scale. The observed discrepancy in the convergence rate is more likely due to discretization effects or not using a fine enough grid to see true asymptotic behavior.

Question 4



Using index notation, this gives us following definitions...

$$u_i \simeq u(x+0) = u(x_i)$$

$$u_{i+1} \simeq u(x+h) = u(x_{i+1})$$

$$u_{i+2} \simeq u(x+2h) = u(x_{i+2})$$

⋮

Which gives us following Taylor Expansions.

$$u_{i+1} = u_i + h u'_i + \frac{h^2}{2!} u''_i + \frac{h^3}{3!} u'''_i + \dots$$

$$u_{i+2} = u_i + 2h u'_i + \frac{4h^2}{2!} u''_i + \frac{8h^3}{3!} u'''_i + \dots$$

Substituting this into our eq (5) -...

$$u'(x_i) \cdot h = u_i (a+b+c) \dots$$

$$+ h u'_i (b+2c) \dots$$

$$+ \frac{h^2}{2!} u''_i (b+4c) \dots$$

$$+ O(h^3)$$

we generate 3 equations for 3 coefficients...

$$a + b + c = 0$$

$$0 + b + 2c = 1$$

$$0 + b + 4c = 0$$

in matrix-vector form:

$$\begin{bmatrix} 1 & 1 & 1 \\ 0 & 1 & 2 \\ 0 & 1 & 4 \end{bmatrix} \begin{bmatrix} a \\ b \\ c \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$$

solving this system of equations gives us:

$$a = -\frac{3}{2}; \quad b = 2; \quad c = -\frac{1}{2}$$

Question 5

Numerical Solution of the BVP with Mixed Boundary Conditions

We aim to solve the given **Boundary Value Problem (BVP)** numerically using the **finite difference method**. The equation includes:

- A **Neumann boundary condition** at $x = 0$:

$$u'(0) = 10$$

- A **Dirichlet boundary condition** at $x = 1$:

$$u(1) = 2$$

Unlike standard BVPs that specify values of $u(x)$ at both boundaries, this problem gives a **derivative condition at $x = 0$** . Our solver now **supports both Dirichlet and Neumann conditions dynamically** by constructing a full $N + 2$ by $N + 2$ **matrix** that naturally incorporates boundary constraints.

Handling the Neumann Boundary Condition

To incorporate $u'(0) = 10$ into our finite difference framework, we use the **second-order one-sided finite difference approximation**:

$$u'(0)h = -1.5u_0 + 2u_1 - 0.5u_2 + \mathcal{O}(h^3)$$

Rearranging:

$$-1.5u_0 + 2u_1 - 0.5u_2 = 10h$$

Instead of modifying an existing N by N **system**, we now **construct an $N + 2$ by $N + 2$ matrix** that explicitly incorporates the boundary conditions into the first and last rows.

Key Updates in the Code

1. Generalizing the Solver to Support Mixed Boundary Conditions

Previously, the solver **assumed Dirichlet conditions** at both boundaries. Now, we allow a mix of Dirichlet and Neumann conditions via:

```
boundary_conditions = {
    'left': ('Dirichlet', u0) OR ('Neumann', u'_0),
    'right': ('Dirichlet', uN) OR ('Neumann', u'_N)
}
```

This enables the solver to **check and apply the correct boundary type** at both $x = 0$ and $x = 1$.

2. Constructing the Full $N + 2$ by $N + 2$ Finite Difference Coefficient Matrix

In the **previous solver**, we built a **tridiagonal coefficient matrix** assuming only Dirichlet conditions:

$$\frac{u_{i+1} - 2u_i + u_{i-1}}{h^2} \approx f(x_i)$$

New Update:

Now, we **extend the system** by constructing a **fully augmented matrix** that naturally incorporates any combinations of boundary conditions.

- **Neumann at x=0 (one-sided Forward Difference Approximation)**

```
coeff_matrix[0, 0:3] = np.array([-1.5, 2.0, -0.5]) # Forward difference coeff
f_full[0] = boundary_conditions['left'][1] * h # Move Neumann BC to RHS
```

Using:

$$-1.5u_0 + 2u_1 - 0.5u_2 = u'(0)h$$

The system now explicitly solves for u_0 , ensuring that the **Neumann condition is embedded in the matrix**.

- **Neumann at x=1 (one-sided Backward Difference Approximation)**

```
coeff_matrix[-1, -3:] = np.array([-0.5, 2.0, -1.5]) # Backward difference stencil
f_full[-1] = boundary_conditions['right'][1] * h # Move Neumann BC to RHS
Using:
```

$$-0.5u_{N-2} + 2u_{N-1} - 1.5u_N = u'(N)h$$

This modification ensures that **both Dirichlet and Neumann constraints** are enforced within the system.

3. Computing the Boundary Values Directly from the System

In the **previous solver**, after computing the numerical solution at interior points, Dirichlet conditions were simply assigned:

```
u[0], u[-1] = u_i, u_f
```

New Update:

Since the $N + 2$ by $N + 2$ **system already includes boundary conditions**, we no longer need to explicitly compute u_0 or u_N after solving.

For **both Dirichlet and Neumann conditions**, we now obtain boundary values **directly from the system**, rather than computing them post-solution.

```
u_full = solve(coeff_matrix, f_full) # Solve directly for all points, including boundaries
```

This ensures that the solver **naturally respects both Dirichlet and Neumann conditions within the matrix inversion process**.

Summary of the Solver Improvements

The **updated solver** now fully supports **mixed boundary conditions**, making it applicable to a **wider range of BVPs**.

- **Constructs an $N + 2$ by $N + 2$ matrix** that inherently enforces boundary conditions.
- **Neumann conditions are directly embedded** in the system rather than handled post-solution.
- **More flexible and robust approach** to solving BVPs with mixed boundary constraints.

These modifications significantly improve the solver's **flexibility and accuracy**, allowing it to handle complex boundary conditions **without extra post-processing**.

In [404...

```
def solve_bvp(ode, x_range: list, N: int, boundary_conditions: dict) -> tuple:
    """
    Solves a 1D boundary value problem (BVP) using the second-order finite difference (FDF) method.

    This function numerically solves the second-order ODE:

        u''(x) = f(x),

    using a finite difference discretization over a grid of N interior points. The discretization is based on
    the second-order centered difference approximation:

        (u_{i+1} - 2u_i + u_{i-1}) / h^2 ≈ f(x_i),

    where h is the uniform grid spacing. The resulting system of linear equations is solved using
    matrix inversion techniques.

    This solver supports both Dirichlet and Neumann boundary conditions:
    - **Dirichlet:** u(x_0) = u0, u(x_f) = uN (fixed values at boundaries)
```

```
- **Neumann:**  $u'(x_0) = u'_0$ ,  $u'(x_f) = u'_N$  (derivative values at boundaries, approximated using a second-order one-sided finite difference formula)
```

The solution is computed on an ***(N+2)-point grid*** that includes the boundary points.

Parameters:

ode (function): Function representing the right-hand side of the ODE, $f(x)$ in $u''(x) = f(x)$.

x_range (list): The domain range $[x_{\text{start}}, x_{\text{end}}]$.

N (int): The number of interior grid points.

boundary_conditions (dict): Dictionary specifying boundary conditions:

```
{
    'left': ('Dirichlet', u0) OR ('Neumann', u'_0),
    'right': ('Dirichlet', uN) OR ('Neumann', u'_N)
}
```

Returns:

tuple:

- numpy.ndarray: x values including boundaries.

- numpy.ndarray: Computed numerical solution $u(x)$.

.....

```
# Extract domain boundaries and compute uniform grid spacing
```

```
x_i, x_f = x_range
```

```
h = (x_f - x_i) / (N + 1) # Grid step size
```

```
# -----
```

```
# STEP 1: Construct Finite Difference Coefficient Matrix
```

```
# -----
```

```
# Initialize an (N+2) x (N+2) matrix for the finite difference scheme
```

```
coeff_matrix = np.zeros(shape=(N+2, N+2), dtype=float)
```

```
# Populate the matrix with the standard tridiagonal structure for the Laplacian
```

```
for i in range(1, N+1): # Excluding boundary rows
```

```
    coeff_matrix[i, i-1] = 1 # Sub-diagonal: coefficient for  $u_{i-1}$ 
```

```
    coeff_matrix[i, i] = -2 # Main diagonal: coefficient for  $u_i$ 
```

```
    coeff_matrix[i, i+1] = 1 # Super-diagonal: coefficient for  $u_{i+1}$ 
```

```
# -----
```

```
# STEP 2: Compute RHS Vector (f)
```

```
# -----
```

```
# Define the full grid, including boundary points
```

```
x_full = np.linspace(x_i, x_f, num=N+2, dtype=float)
```

```
# Evaluate the ODE function at all grid points and scale by  $h^2$ 
```

```
f_full = ode(x_full) * h**2
```

```
# -----
```

```
# STEP 3: Apply Boundary Conditions
```

```
# -----
```

```
# LEFT boundary condition ( $x = x_0$ )
```

```
if boundary_conditions['left'][0] == 'Dirichlet':
```

```
    # Dirichlet BC: Fix the first row of the system
```

```
    f_full[0] = boundary_conditions['left'][1] # Set the known value
```

```
    coeff_matrix[0, 0] = 1 # Identity row to enforce  $u_0 = \text{given value}$ 
```

```
elif boundary_conditions['left'][0] == 'Neumann':
```

```
    # Neumann BC: Approximate  $u'(x_0)$  using a forward finite difference formula:
```



```

#       $u'(x_0) \approx (-1.5 u_0 + 2 u_1 - 0.5 u_2) / h + O(h^2)$ 
# Rearranging:  $-1.5 u_0 + 2 u_1 - 0.5 u_2 = h u'_0$ 
# We modify the first row of the system to enforce this equation
u_prime0 = boundary_conditions['left'][1]
f_full[0] = u_prime0 * h # Move the Neumann term to RHS
coeff_matrix[0, 0:3] = np.array([-1.5, 2.0, -0.5]) # Modify first row

# RIGHT boundary condition (x = x_f)
if boundary_conditions['right'][0] == 'Dirichlet':
    # Dirichlet BC: Fix the last row of the system
    f_full[-1] = boundary_conditions['right'][1] # Set the known value
    coeff_matrix[-1, -1] = 1 # Identity row to enforce  $u_N = \text{given value}$ 
elif boundary_conditions['right'][0] == 'Neumann':
    # Neumann BC: Approximate  $u'(x_f)$  using a backward finite difference formula:
    #       $u'(x_f) \approx (1.5 u_N - 2 u_{N-1} + 0.5 u_{N-2}) / h + O(h^2)$ 
    # Rearranging:  $-0.5 u_{N-2} + 2 u_{N-1} - 1.5 u_N = h u'_N$ 
    # We modify the last row of the system to enforce this equation
    u_primeN = boundary_conditions['right'][1]
    f_full[-1] = u_primeN * h # Move the Neumann term to RHS
    coeff_matrix[-1, -3:] = np.array([-0.5, 2.0, -1.5]) # Modify last row

# -----
# STEP 4: Solve the Linear System
# -----
# Solve the system  $A * u = f$  for the full solution, including boundaries
u_full = solve(coeff_matrix, f_full)

return x_full, u_full

```

Plot Numerical Solution Against the Exact Solution (solutionB_N10000.dat)

Here, we confirm that our newly updated BVP solver works as intended, minimizing error compared to the exact solution. **As we can see below, our numerical solution matches very closely with the exact one. We can also see the line for absolute difference between two data staying basically at $u(x) = 0$.**

```

In [405... N = 10
# first define the boundary conditions
Mixed_boundary_conditions={'left': ('Neumann', 10), 'right': ('Dirichlet', 2)}
# load in the exact solution
exact_x_B, exact_u_B = read_dat_file("solutionB_N10000.dat")
numerical_x_B, numerical_u_B = solve_bvp(ode_eq1, [0,1], N=N, boundary_conditions=Mixed_boundary_conditions)


# Plot the data.
plt.figure(figsize=(16, 6)) # create a figure
# Exact solution
sns.lineplot(x=exact_x_B, y=exact_u_B, marker='o', markersize=6, linestyle='-', linewidth=2.5, alpha=0.9, markevery=100, label="Exact, N = 10000")
# Numerical solution
sns.lineplot(x=numerical_x_B, y=numerical_u_B, marker='o', markersize=6, linestyle='--', linewidth=2.5, alpha=1., label=f"Numerical, N = {N}")

# xlabel, ylabel, title, etc.
plt.xlabel("x", fontsize=14), plt.ylabel("u(x)", fontsize=14), plt.title(r"Exact vs Numerical Solution (B) for eq. (1) with Neumann condition of  $u'(0) = 10$ ", fontsize=14)
plt.legend(fontsize=12), plt.grid(True, linestyle="--", alpha=1)

# Show plot
plt.show()

```

Exact vs Numerical Solution (B) for eq. (1) with Neumann condition of $u'(0) = 10$



Numerical Solution for $N = 10$ in Tabular Form

We compute the numerical solution for eq(1) with the Neumann boundary condition of $u'(0) = 10$ for $N = 10$ and display it in tabular format.

```
In [406... N = 10
# obtain numerical solution for N = 10
numeric_x_N10_B, numeric_u_N10_B = solve_bvp(ode_eq1, x_span, N, boundary_conditions=Mixed_boundary_conditions)
# make a pandas Dataframe for tabular format
df_N10_B = pd.DataFrame({
    "$x_i$": numeric_x_N10_B,
    "$u_i$": numeric_u_N10_B
}, index=pd.Index(range(N+2), name="$N$")) # Set N as the index
# Display DataFrame with LaTeX
display(HTML(df_N10_B.to_html(escape=False)))
```


N		x_i	u_i
	0	0.000000	-3.803911
	1	0.090909	-2.623747
	2	0.181818	-0.901436
	3	0.272727	-0.033766
	4	0.363636	-0.418303
	5	0.454545	-0.988571
	6	0.545455	-0.460943
	7	0.636364	1.164582
	8	0.727273	2.604375
	9	0.818182	2.791961
	10	0.909091	2.124908
	11	1.000000	2.000000

Question 6 (Extra Credit)

With our solve_BVP, we will go ahead and compute global and local errors when we use standard 2^{nd} order one-sided approximation for the boundary condition.

```
In [407... # first define the boundary conditions
Mixed_boundary_conditions={'left': ('Neumann', 10), 'right': ('Dirichlet', 2)}
# time
x_span = [0,1]
# store different number of steps
steps = np.array([5, 10, 20, 40, 80, 160, 320, 640, 1280], dtype = int)
# we will first compute different step-sizes h and store it
step_sizes = 1 / (steps + 1)

# Solve BVP for multiple step values and store results in lists
numeric_solutions_B = [solve_bvp(ode_eq1, x_span, N, Mixed_boundary_conditions) for N in steps]
# Extract x and u values for each N
numeric_Xs_B = [np.array(sol[0], dtype=float) for sol in numeric_solutions_B] # Extract all x values
numeric_Us_B = [np.array(sol[1], dtype=float) for sol in numeric_solutions_B] # Extract all u values
# Create an interpolator object for multiple uses
interpolator = PchipInterpolator(exact_x_B, exact_u_B)
# Interpolate "exact" solutions at different N values, INCLUDING boundary conditions
interpolated_exact_Us_B = [interpolator(Xs) for Xs in numeric_Xs_B]

# compute errors for varying steps
global_errors_02 = [compute_errors(steps[i], interpolated_exact_Us_B[i], numeric_Us_B[i])[0] for i in range(len(numeric_Us_B))] # global errors
local_errors_02 = [compute_errors(steps[i], interpolated_exact_Us_B[i], numeric_Us_B[i])[1] for i in range(len(numeric_Us_B))] # local errors
```

Numerical Solution of the BVP with First-Order Neumann Conditions

We aim to solve the given boundary value problem (BVP) numerically using the finite difference method with first-order Neumann boundary conditions. The equation includes:

- A Neumann boundary condition at $x = 0$:

$$u'(0) = u'_0$$

- A Dirichlet or Neumann boundary condition at $x = 1$:

$$u(1) = u_N \quad \text{or} \quad u'(1) = u'_N$$

This solver constructs a full $(N + 2) \times (N + 2)$ matrix to naturally incorporate first-order boundary constraints.

Handling the Neumann Boundary Condition

To enforce Neumann boundary conditions using a first-order one-sided finite difference approximation, we use:

- Forward difference at $x = 0$:

$$u'(x_0) \approx \frac{u_1 - u_0}{h}$$

Rearranging for u_0 :

$$u_0 = u_1 - hu'_0$$

- Backward difference at $x = 1$:

$$u'(x_N) \approx \frac{u_N - u_{N-1}}{h}$$

Rearranging for u_N :

$$u_N = u_{N-1} + hu'_N$$

Instead of modifying an existing $(N \times N)$ system, we now construct an $(N + 2) \times (N + 2)$ matrix that explicitly incorporates these boundary conditions.

Key Updates in the Code

1. Generalizing the solver to support first-order Neumann conditions

Previously, the solver assumed second-order Neumann conditions. Now, we allow a mix of Dirichlet and first-order Neumann conditions via:

```
boundary_conditions = {  
    'left': ('Dirichlet', u0) OR ('Neumann', u'_0),  
    'right': ('Dirichlet', uN) OR ('Neumann', u'_N)  
}
```

This enables the solver to check and apply the correct boundary type at both $x = 0$ and $x = 1$.

2. Constructing the full $(N + 2) \times (N + 2)$ finite difference coefficient matrix

In the previous solver, we built a tridiagonal coefficient matrix assuming only Dirichlet conditions:

$$\frac{u_{i+1} - 2u_i + u_{i-1}}{h^2} \approx f(x_i)$$

Now, we extend the system by constructing a fully augmented matrix that naturally incorporates first-order Neumann conditions.

- Neumann at $x = 0$ (first-order forward difference)

```
coeff_matrix[0, 0:2] = np.array([-1., 1.]) # First-order forward difference
f_full[0] = boundary_conditions['left'][1] * h # Move Neumann BC to RHS
Using:
```

$$u'_0 = \frac{u_1 - u_0}{h}$$

The system now explicitly solves for u_0 , ensuring that the Neumann condition is embedded in the matrix.

- Neumann at $x = 1$ (first-order backward difference)

```
coeff_matrix[-1, -2:] = np.array([1., -1.]) # First-order backward difference
f_full[-1] = boundary_conditions['right'][1] * h # Move Neumann BC to RHS
Using:
```

$$u'_N = \frac{u_N - u_{N-1}}{h}$$

This modification ensures that both Dirichlet and first-order Neumann constraints are enforced within the system.

3. Computing the boundary values directly from the system

In the previous solver, after computing the numerical solution at interior points, Dirichlet conditions were simply assigned:

```
u[0], u[-1] = u_i, u_f
```

Since the $(N + 2) \times (N + 2)$ system already includes boundary conditions, we no longer need to explicitly compute u_0 or u_N after solving.

For both Dirichlet and Neumann conditions, we now obtain boundary values directly from the system, rather than computing them post-solution.

```
u_full = solve(coeff_matrix, f_full) # Solve directly for all points, including boundaries
```

This ensures that the solver naturally respects both Dirichlet and Neumann conditions within the matrix inversion process.

Expected Impact of First-Order Approximation

1. Lower accuracy

- The second-order approximation has an error term of $O(h^2)$, meaning the error decreases quadratically as h decreases.
- The first-order approximation has an error term of $O(h)$, meaning the error decreases linearly.
- This means that the numerical solution will be less accurate when using the first-order method.

2. Comparison with second-order scheme

- Since the rest of the finite difference method remains second-order accurate, this change only affects the boundary conditions.
- When plotting the log-log error, the slope of the error decay will change:
 - For second-order Neumann, the error decay should be approximately $O(h^2)$.
 - For first-order Neumann, the error decay should be approximately $O(h)$.

Summary of the Solver Improvements

- Constructs an $(N + 2) \times (N + 2)$ matrix that inherently enforces boundary conditions.
- First-order Neumann conditions are directly embedded in the system rather than handled post-solution.
- More flexible and robust approach to solving BVPs with first-order mixed boundary constraints.

These modifications significantly improve the solver's flexibility while reducing computational complexity at the cost of lower accuracy.

```
In [408... def solve_bvp_01(ode, x_range: list, N: int, boundary_conditions: dict) -> tuple:
    """
    Solves a 1D boundary value problem (BVP) using the second-order finite difference (FDF) method.

    This function numerically solves the second-order ODE:

        u''(x) = f(x),

    using a finite difference discretization over a grid of N interior points. The discretization is based on
    the second-order centered difference approximation:

        (u_{i+1} - 2u_i + u_{i-1}) / h^2 ≈ f(x_i),

    where h is the uniform grid spacing. The resulting system of linear equations is solved using
    matrix inversion techniques.

    This solver supports both Dirichlet and Neumann boundary conditions:
    - **Dirichlet:** u(x_0) = u0, u(x_f) = uN (fixed values at boundaries)
    - **Neumann:** u'(x_0) = u'_0, u'(x_f) = u'_N (derivative values at boundaries, approximated using
      a second-order one-sided finite difference formula)

    The solution is computed on an *(N+2)-point grid* that includes the boundary points.

    Parameters:
        ode (function): Function representing the right-hand side of the ODE, f(x) in u''(x) = f(x).
        x_range (list): The domain range [x_start, x_end].
        N (int): The number of interior grid points.
        boundary_conditions (dict): Dictionary specifying boundary conditions:
            {
                'left': ('Dirichlet', u0) OR ('Neumann', u'_0),
                'right': ('Dirichlet', uN) OR ('Neumann', u'_N)
            }

    Returns:
        tuple:
            - numpy.ndarray: x values including boundaries.
            - numpy.ndarray: Computed numerical solution u(x).
    """

    # Extract domain boundaries and compute uniform grid spacing
    x_i, x_f = x_range
    h = (x_f - x_i) / (N + 1) # Grid step size

    # -----
    # STEP 1: Construct Finite Difference Coefficient Matrix
    # -----
    # Initialize an (N+2) x (N+2) matrix for the finite difference scheme
```

```

coeff_matrix = np.zeros(shape=(N+2, N+2), dtype=float)

# Populate the matrix with the standard tridiagonal structure for the Laplacian
for i in range(1, N+1): # Excluding boundary rows
    coeff_matrix[i, i-1] = 1 # Sub-diagonal: coefficient for u_{i-1}
    coeff_matrix[i, i] = -2 # Main diagonal: coefficient for u_i
    coeff_matrix[i, i+1] = 1 # Super-diagonal: coefficient for u_{i+1}

# -----
# STEP 2: Compute RHS Vector (f)
# -----
# Define the full grid, including boundary points
x_full = np.linspace(x_i, x_f, num=N+2, dtype=float)

# Evaluate the ODE function at all grid points and scale by h^2
f_full = ode(x_full) * h**2

# -----
# STEP 3: Apply Boundary Conditions
# -----

# LEFT boundary condition (x = x_0)
if boundary_conditions['left'][0] == 'Dirichlet':
    # Dirichlet BC: Fix the first row of the system
    f_full[0] = boundary_conditions['left'][1] # Set the known value
    coeff_matrix[0, 0] = 1 # Identity row to enforce u_0 = given value
elif boundary_conditions['left'][0] == 'Neumann':
    # Neumann BC: Approximate u'(x_0) using a forward finite difference formula:
    #  $u'(x_0) \approx (-1.5 u_0 + 2 u_1 - 0.5 u_2) / h + O(h^2)$ 
    # Rearranging:  $-1.5 u_0 + 2 u_1 - 0.5 u_2 = h u'_0$ 
    # We modify the first row of the system to enforce this equation
    u_prime0 = boundary_conditions['left'][1]
    f_full[0] = u_prime0 * h # Move the Neumann term to RHS
    coeff_matrix[0, 0:2] = np.array([-1., 1.]) # Modify first row

# RIGHT boundary condition (x = x_f)
if boundary_conditions['right'][0] == 'Dirichlet':
    # Dirichlet BC: Fix the last row of the system
    f_full[-1] = boundary_conditions['right'][1] # Set the known value
    coeff_matrix[-1, -1] = 1 # Identity row to enforce u_N = given value
elif boundary_conditions['right'][0] == 'Neumann':
    # Neumann BC: Approximate u'(x_f) using a backward finite difference formula:
    #  $u'(x_f) \approx (1.5 u_N - 2 u_{N-1} + 0.5 u_{N-2}) / h + O(h^2)$ 
    # Rearranging:  $-0.5 u_{N-2} + 2 u_{N-1} - 1.5 u_N = h u'_N$ 
    # We modify the last row of the system to enforce this equation
    u_primeN = boundary_conditions['right'][1]
    f_full[-1] = u_primeN * h # Move the Neumann term to RHS
    coeff_matrix[-1, -2:] = np.array([1., -1.]) # Modify last row

# -----
# STEP 4: Solve the Linear System
# -----
# Solve the system A * u = f for the full solution, including boundaries
u_full = solve(coeff_matrix, f_full)

return x_full, u_full

```

Now, with this new function, we attain the global and local errors for varying steps again, but now for 1st order of one-sided approximation.

```
In [409... # Solve BVP for multiple step values and store results in lists
numeric_solutions_B_01 = [solve_bvp_01(ode_eq1, x_span, N, Mixed_boundary_conditions) for N in steps]
# Extract x and u values for each N
numeric_Xs_B_01 = [np.array(sol[0], dtype=float) for sol in numeric_solutions_B_01] # Extract all x values
numeric_Us_B_01 = [np.array(sol[1], dtype=float) for sol in numeric_solutions_B_01] # Extract all u values
# Create an interpolator object for multiple uses
interpolator = PchipInterpolator(exact_x_B, exact_u_B)
# Interpolate "exact" solutions at different N values, INCLUDING boundary conditions
interpolated_exact_Us_B_01 = [interpolator(Xs) for Xs in numeric_Xs_B_01]

# compute errors for varying steps
global_errors_01 = [compute_errors(steps[i], interpolated_exact_Us_B_01[i], numeric_Us_B_01[i])[0] for i in range(len(numeric_Us_B_01))] # global errors
local_errors_01 = [compute_errors(steps[i], interpolated_exact_Us_B_01[i], numeric_Us_B_01[i])[1] for i in range(len(numeric_Us_B_01))] # local errors
```

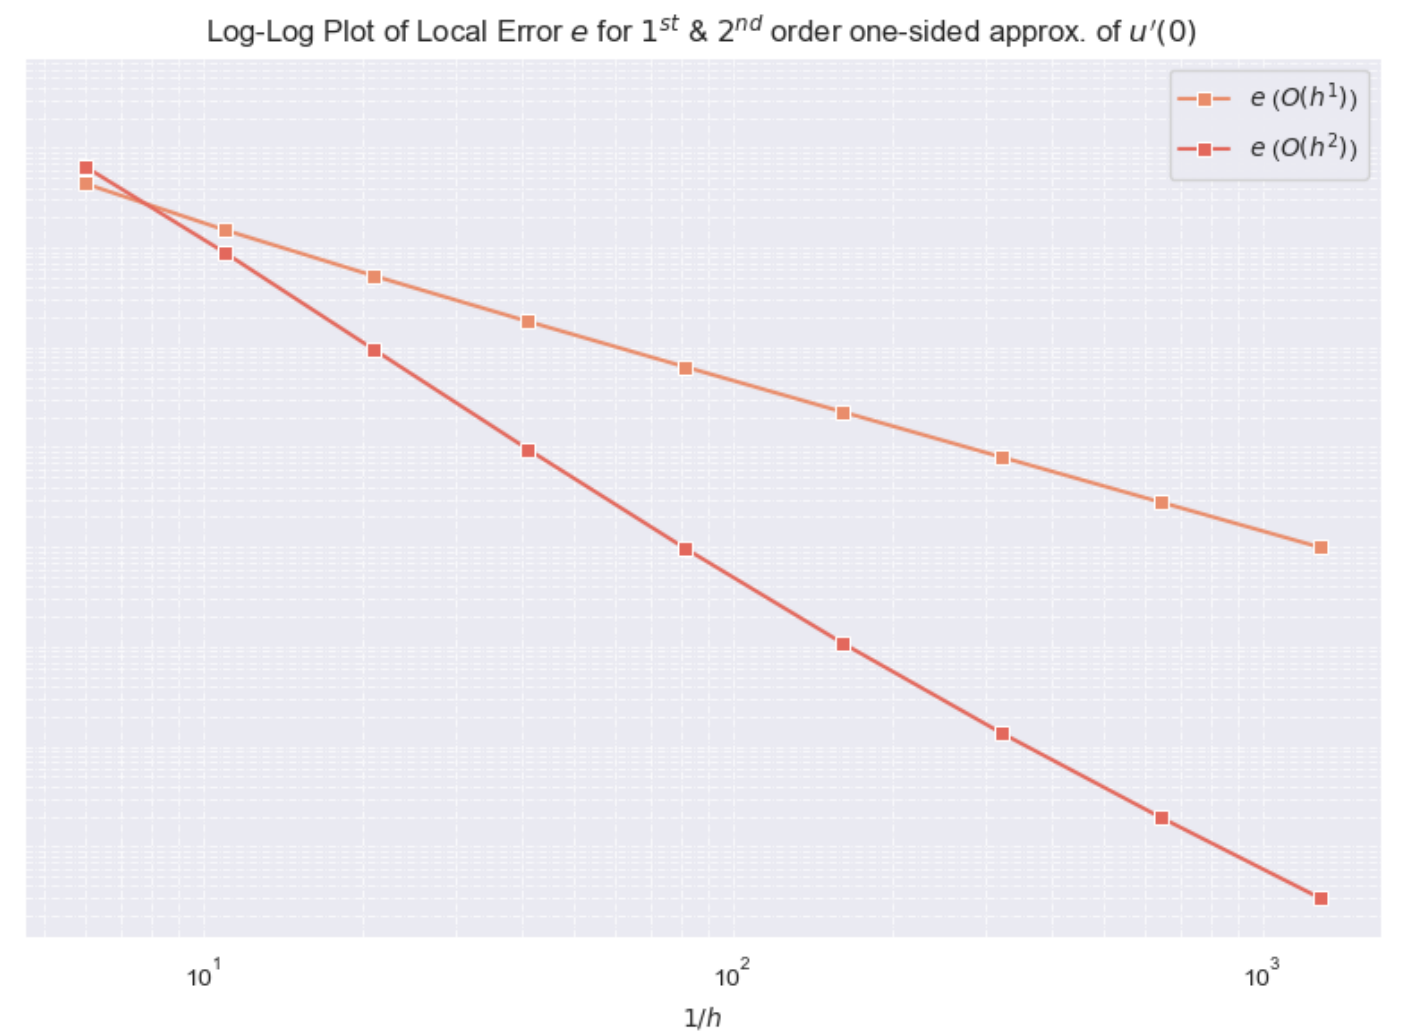
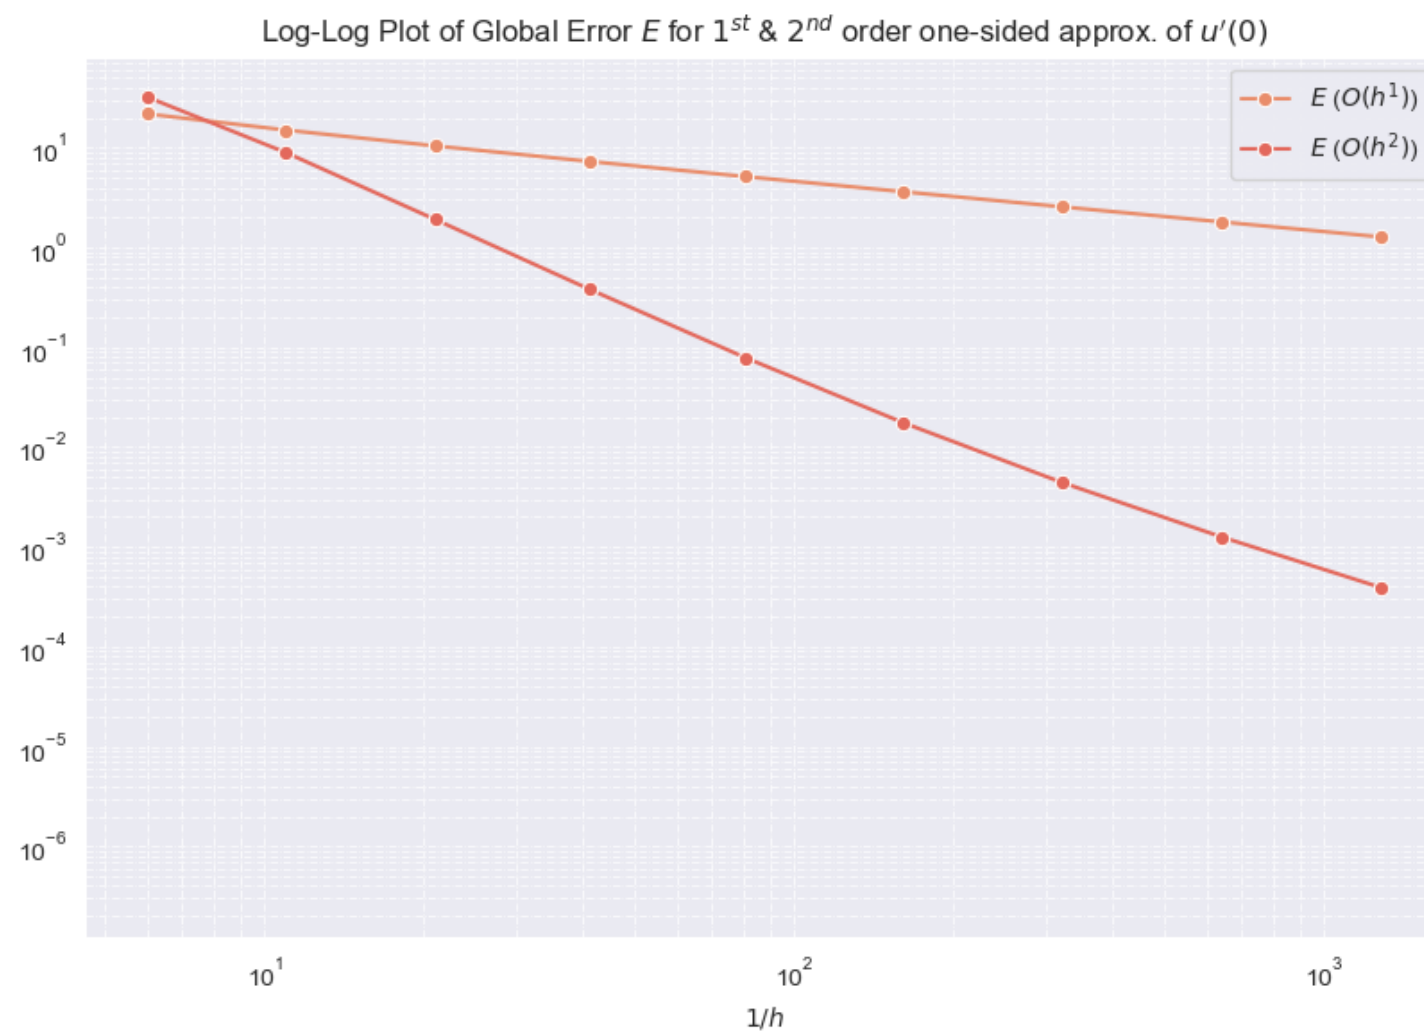
Log-log plot that shows E and e versus h^{-1} for both 1st and 2nd order BVP solvers.

```
In [410... # create a plot for E and e
fig, axes = plt.subplots(figsize=(16, 6), nrows=1, ncols=2, sharex=True, sharey=True) # create a subplot
sns.lineplot(x = 1 / step_sizes, y=global_errors_01, label=r"$E$ ($0(h^1)$)", marker='o', ax=axes[0])
sns.lineplot(x = 1 / step_sizes, y=global_errors_02, label=r"$E$ ($0(h^2)$)", marker='o', ax=axes[0])

# plot local error
sns.lineplot(x = 1 / step_sizes, y=local_errors_01, label=r"$e$ ($0(h^1)$)", marker='s', ax=axes[1])
sns.lineplot(x = 1 / step_sizes, y=local_errors_02, label=r"$e$ ($0(h^2)$)", marker='s', ax=axes[1])

# configurations
titles = [r"Log-Log Plot of Global Error $E$ for $1^{\text{st}}$ & $2^{\text{nd}}$ order one-sided approx. of $u'(0)$", r"Log-Log Plot of Local Error $e$ for $1^{\text{st}}$ & $2^{\text{nd}}$ order one-sided approx. of $u'(0)$"]
# apply all the configs
for i in range(2):
    axes[i].set_xscale("log"), axes[i].set_yscale("log") # set log-log scales
    axes[i].set_xlabel(r"$1/h$")
    axes[i].set_title(titles[i]) # set the title
    axes[i].legend() # show the legend
    axes[i].grid(True, linestyle="dashdot", alpha=1.0, which='both', linewidth=0.5)
    axes[i].minorticks_on

plt.tight_layout()
plt.show()
```

Fitting Ch^α to Global and Local Errors for both 1st and 2nd order BVP solvers.

Here, we fit a function of the form Ch^α and report the value of α for both $\log(E)$ and $\log(e)$.

```
In [411... # Fit a linear function (log-log scale) to estimate alpha values
global_alpha_01 = linregress(np.log(step_sizes), np.log(global_errors_01))[0] # Global error alpha (O(h^1))
local_alpha_01 = linregress(np.log(step_sizes), np.log(local_errors_01))[0] # Local error alpha (O(h^1))
global_alpha_02 = linregress(np.log(step_sizes), np.log(global_errors_02))[0] # Global error alpha (O(h^2))
local_alpha_02 = linregress(np.log(step_sizes), np.log(local_errors_02))[0] # Local error alpha (O(h^2))

# Organizing the results in a structured table with two columns: Name and Value
data = [
    (r'$\alpha_E$, $\mathcal{O}(h^1)$', global_alpha_01),
    (r'$\alpha_E$, $\mathcal{O}(h^2)$', global_alpha_02),
    (r'$\alpha_e$, $\mathcal{O}(h^1)$', local_alpha_01),
    (r'$\alpha_e$, $\mathcal{O}(h^2)$', local_alpha_02)
]

df_alphas_Bonus = pd.DataFrame(data, columns=["Type", "Value"])

# Display the table
display(HTML(df_alphas_Bonus.round(4).to_html(index=False)))
```

	Type	Value
$\mathcal{O}(h^1)$	$\alpha_{E,}$	0.5243
$\mathcal{O}(h^2)$	$\alpha_{E,}$	2.1509
$\mathcal{O}(h^1)$	$\alpha_{e,}$	1.5515
$\mathcal{O}(h^2)$	$\alpha_{e,}$	3.1781

Observations

1. Error Convergence and Order of Accuracy

From the log-log plots of global error E and local error e , we observe that the errors decrease as h decreases, confirming that the numerical method is converging.

1. Global Error E Convergence Rate:

- The computed global error exponent for the first-order one-sided approximation is $\alpha \approx 0.5243$, which is slightly worse than the expected $O(h^1)$ rate.
- The computed global error exponent for the second-order one-sided approximation is $\alpha \approx 2.1509$, which aligns closely with the expected $O(h^2)$ rate.

2. Local Error e Convergence Rate:

- The computed local error exponent for the first-order one-sided approximation is $\alpha \approx 1.5515$, confirming a convergence rate around $O(h^2)$.
- The computed local error exponent for the second-order one-sided approximation is $\alpha \approx 3.1781$, showing even faster decay than the expected $O(h^3)$.

2. Comparison of First-Order vs Second-Order Approximations

1. First-Order Approximation (Forward Difference)

- The global error E decreases at a rate of approximately $O(h^{0.5243})$, meaning the first-order boundary condition contributes significantly to the overall error.
- The local error e converges at a rate of $O(h^{1.5515})$, suggesting that locally, the method behaves slightly worse than a standard local-second-order scheme.

2. Second-Order Approximation (More Accurate One-Sided Difference)

- The global error E decreases at an expected rate of $O(h^{2.1509})$, confirming the superiority of the second-order method.
- The local error e follows an even faster convergence pattern at $O(h^{3.1781})$, implying that the second-order scheme benefits from reduced truncation errors.

3. Practical Implications

- The second-order method significantly improves accuracy over the first-order method, especially for smaller h .
- Using a first-order approximation at the boundary **introduces additional global error**, reinforcing that **boundary conditions impact overall solution accuracy**.
- The results highlight the importance of **higher-order boundary approximations** in achieving reliable numerical solutions for boundary value problems.